

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")

from keras.preprocessing import sequence
from keras.utils import np_utils
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Embedding
from keras.layers import LSTM, SimpleRNN, GRU
from keras.datasets import imdb
from keras.utils.np_utils import to_categorical
from sklearn.metrics import (precision_score, recall_score, f1_score,
accuracy_score, mean_squared_error, mean_absolute_error)
from sklearn import metrics
from sklearn.svm import SVC
from sklearn.preprocessing import Normalizer
from keras import callbacks
from keras.callbacks import ModelCheckpoint, EarlyStopping,
ReduceLROnPlateau, CSVLogger
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from imblearn.combine import SMOTETomek
from sklearn.linear_model import LogisticRegression
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import KFold
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

```

```

pip install -U imbalanced-learn

```

```

Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: imbalanced-learn in
/usr/local/lib/python3.7/dist-packages (0.9.0)
Collecting imbalanced-learn
  Using cached imbalanced_learn-0.9.1-py3-none-any.whl (199 kB)
Requirement already satisfied: joblib>=1.0.0 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (3.1.0)
Requirement already satisfied: numpy>=1.14.6 in

```

```

/usr/local/lib/python3.7/dist-packages (from imbalanced-learn)
(1.21.6)
Requirement already satisfied: scikit-learn>=1.0.1 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (1.0.2)
Requirement already satisfied: scipy>=1.1.0 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (1.7.3)

```

```
data= pd.read_csv("bank.csv")
```

```

# check shape of dataset
print("shape of the data:", data.shape)
data.head()

```

```
shape of the data: (32950, 16)
```

	age	job	marital	education	default	housing
loan \						
0	49	blue-collar	married	basic.9y	unknown	no
1	37	entrepreneur	married	university.degree	no	no
2	78	retired	married	basic.4y	no	no
3	36	admin.	married	university.degree	no	yes
4	59	retired	divorced	university.degree	no	no

	contact	month	day_of_week	duration	campaign	pdays	previous	\
0	cellular	nov	wed	227	4	999	0	
1	telephone	nov	wed	202	2	999	1	
2	cellular	jul	mon	1148	1	999	0	
3	telephone	may	mon	120	2	999	0	
4	cellular	jun	tue	368	2	999	0	

	poutcome	y
0	nonexistent	no
1	failure	no
2	nonexistent	yes
3	nonexistent	no
4	nonexistent	no

```

# check data types of all columns
data.dtypes

```

age	int64
job	object
marital	object
education	object
default	object
housing	object

```
loan            object
contact         object
month           object
day_of_week     object
duration        int64
campaign        int64
pdays          int64
previous        int64
poutcome        object
y               object
dtype: object
```

```
data.isnull().sum()
```


```
age            0
job            0
marital        0
education      0
default        0
housing        0
loan           0
contact        0
month          0
day_of_week    0
duration       0
campaign       0
pdays         0
previous       0
poutcome       0
y              0
dtype: int64
```

```
# target class count
data["y"].value_counts()
```

```
no      29238
yes      3712
Name: y, dtype: int64
```

```
sns.countplot(data["y"])
plt.title("target variable")
```

```
Text(0.5, 1.0, 'target variable')
```



```
# percentage of class present in target variable(y)
print("percentage of NO and YES\
n",data["y"].value_counts()/len(data)*100)
```


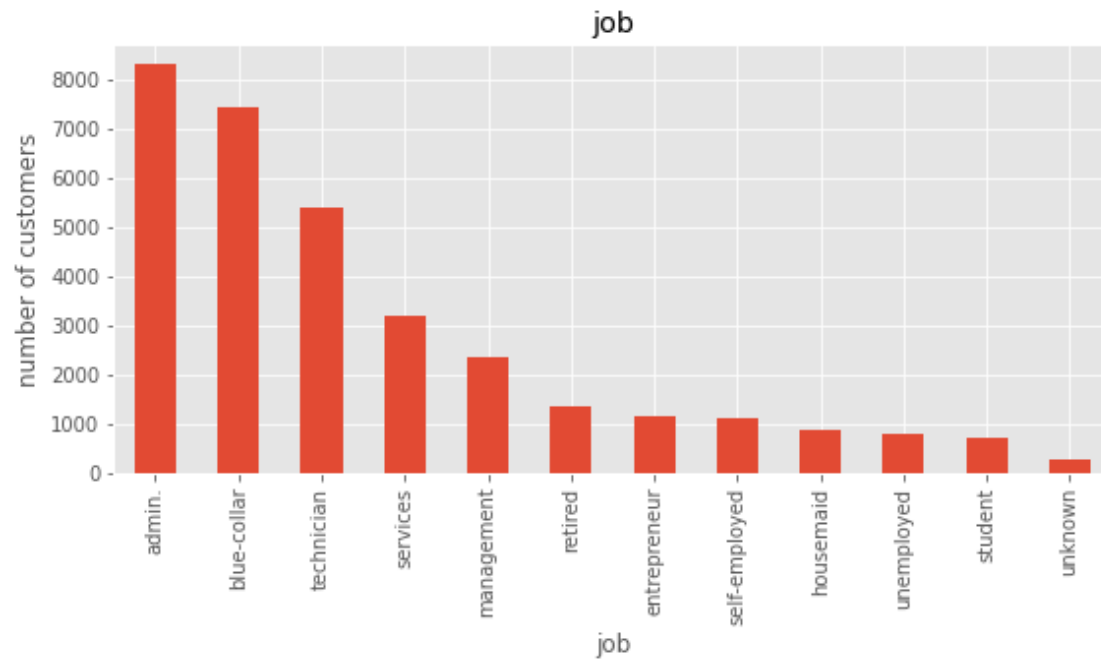
```
percentage of NO and YES
no      88.734446
yes     11.265554
Name: y, dtype: float64
```



```
# indentifying the categorical variables
cat_var= data.select_dtypes(include= ["object"]).columns
print(cat_var)
```

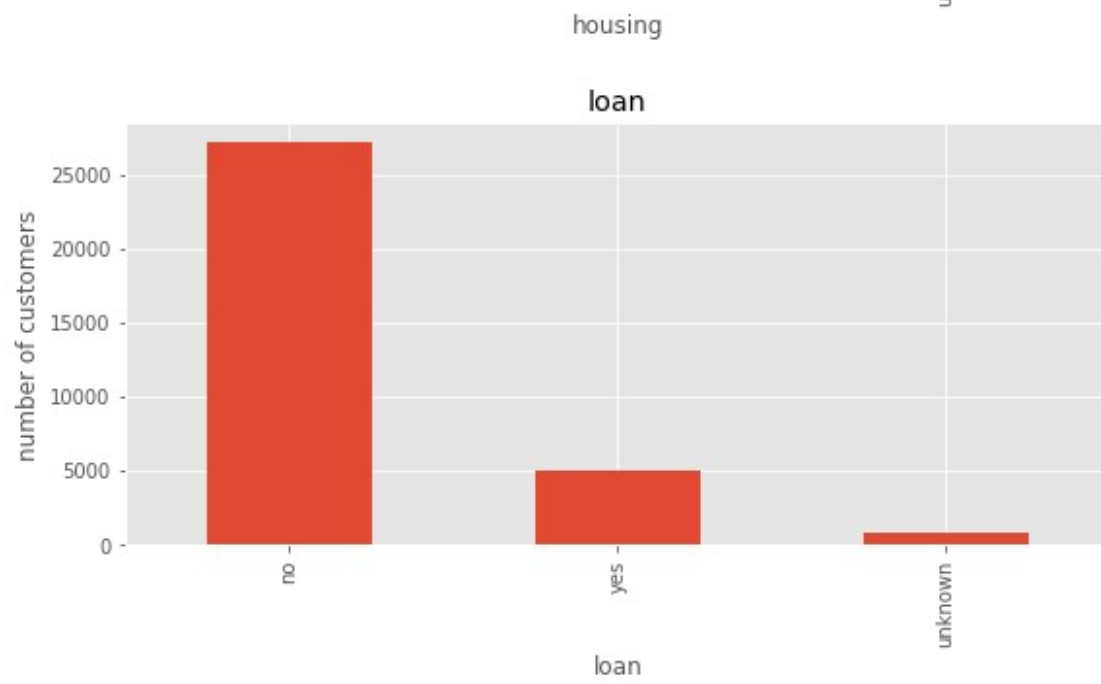

```
# plotting bar chart for each categorical variable
plt.style.use("ggplot")
```



```
for column in cat_var:
    plt.figure(figsize=(20,4))
    plt.subplot(121)
    data[column].value_counts().plot(kind="bar")
    plt.xlabel(column)
    plt.ylabel("number of customers")
    plt.title(column)
```



```
Index(['job', 'marital', 'education', 'default', 'housing', 'loan',
      'contact',
      'month', 'day_of_week', 'poutcome', 'y'],
      dtype='object')
```












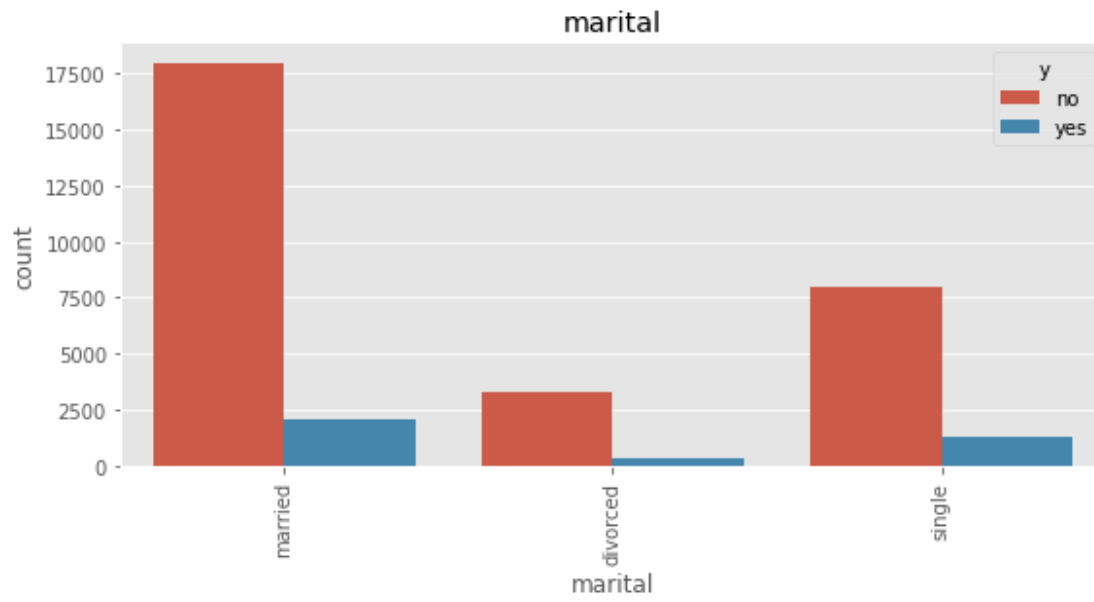








```
# replacing "unknown" with the mode
for column in cat_var:
    mode= data[column].mode()[0]
    data[column]= data[column].replace("unknown", mode)



plt.style.use("ggplot")
for column in cat_var:
    plt.figure(figsize=(20,4))
    plt.subplot(121)
    sns.countplot(data[column], hue=data["y"])
    plt.title(column)
    plt.xticks(rotation=90)
```

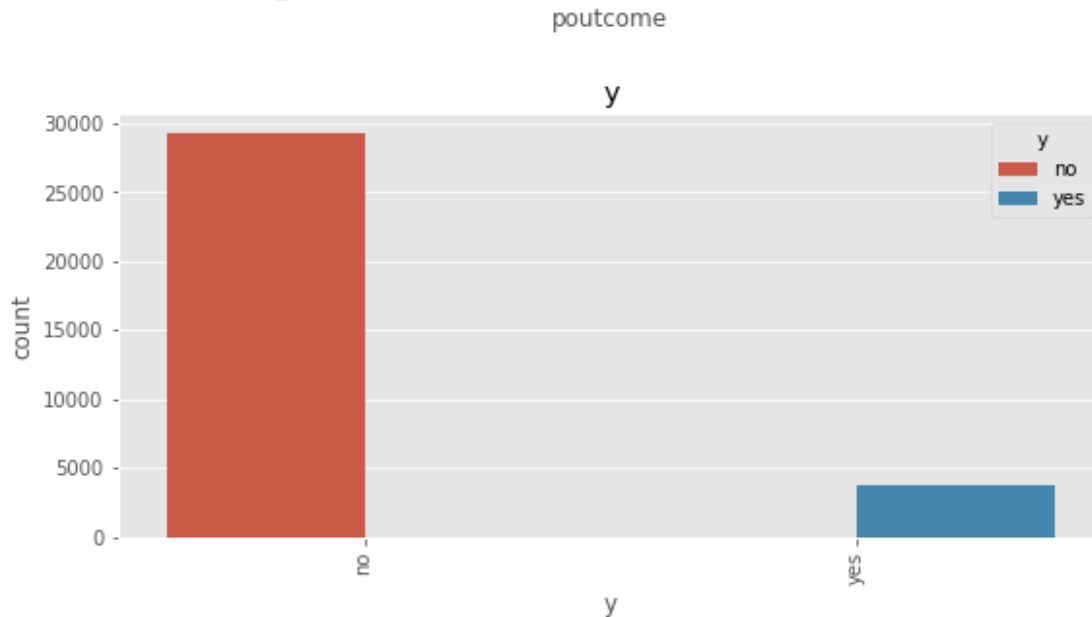
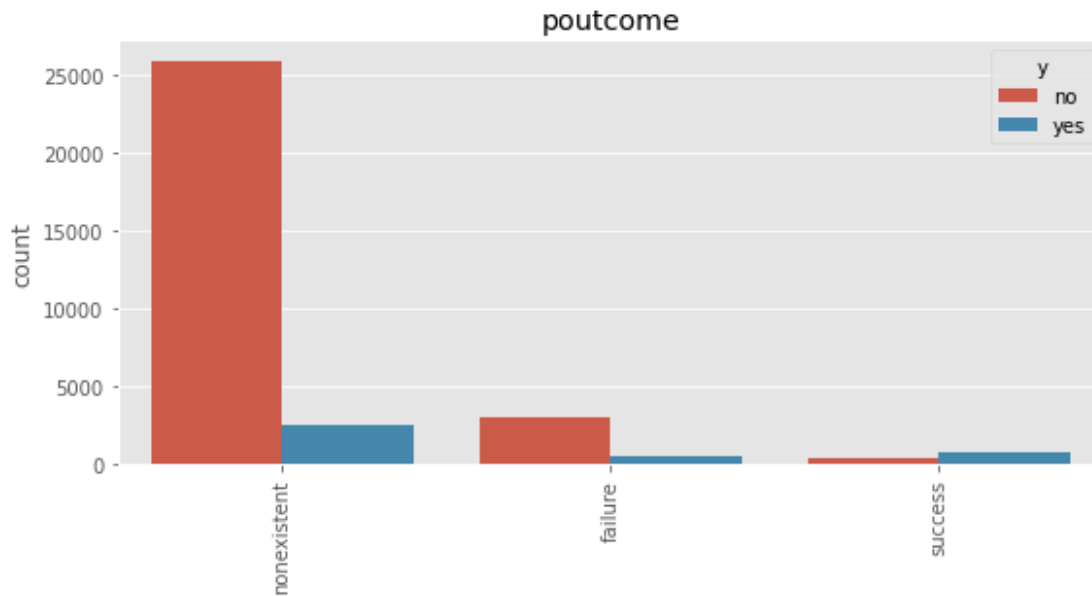












indentifying the numerical variables

```
num_var= data.select_dtypes(include=np.number)
num_var.head()
```



	age	duration	campaign	pdays	previous
0	49	227	4	999	0
1	37	202	2	999	1
2	78	1148	1	999	0
3	36	120	2	999	0
4	59	368	2	999	0

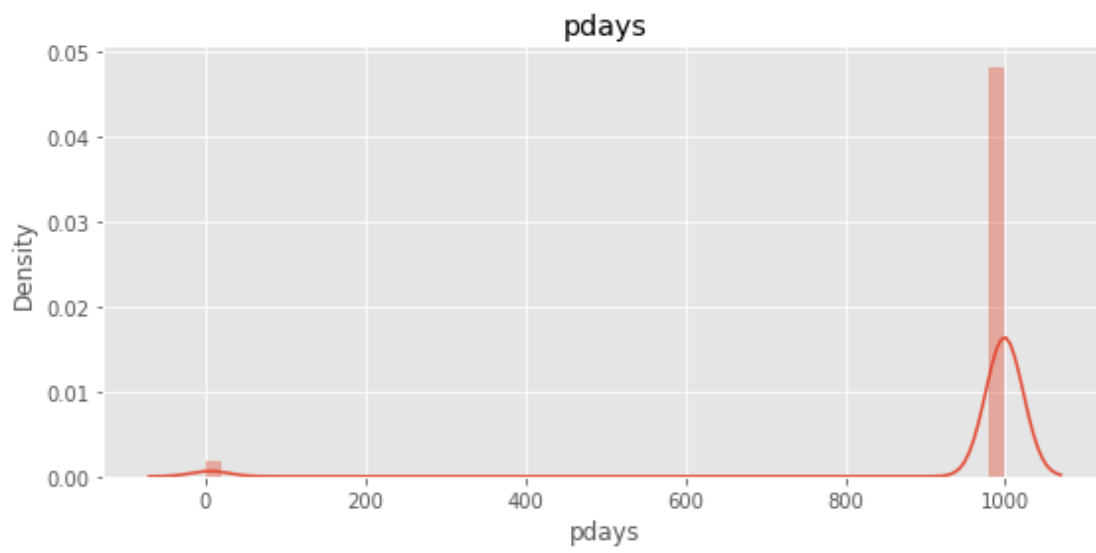

plotting histogram for each numerical variable

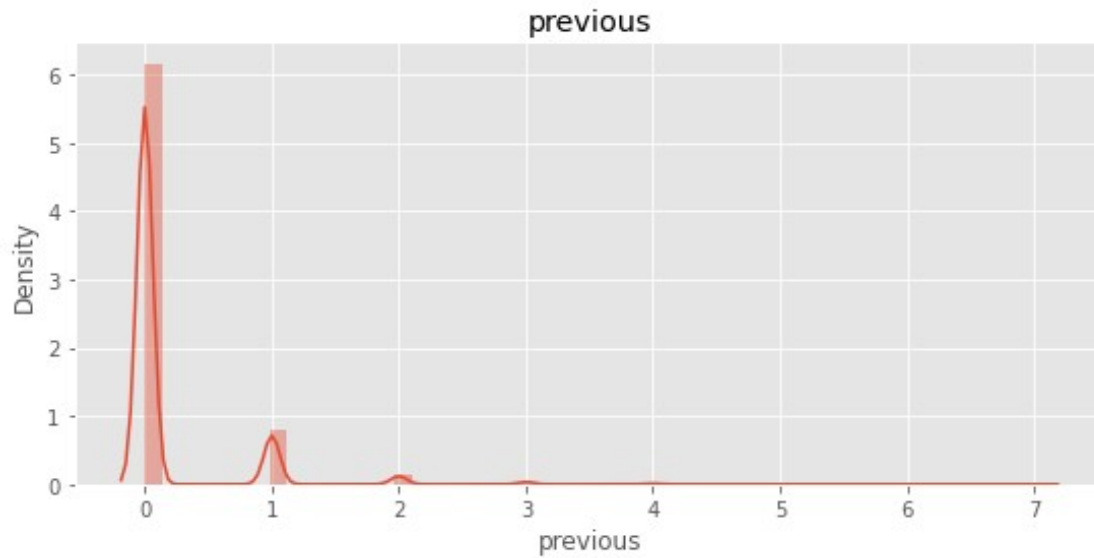
```
plt.style.use("ggplot")
```

```
for column in ["age", "duration", "campaign", "pdays", "previous"]:
```


```
plt.figure(figsize=(20,4))  
plt.subplot(121)  
sns.distplot(data[column], kde=True)  
plt.title(column)
```







```
age = data["age"]  
  
fig = plt.figure(figsize =(10, 7))  
  
# Creating plot  
plt.boxplot(age)  
  
# show plot  
plt.show()
```




```
campaign = data["campaign"]

fig = plt.figure(figsize =(10, 7))

# Creating plot
plt.boxplot(campaign)

# show plot
plt.show()
```




```
duration = data["duration"]

fig = plt.figure(figsize =(10, 7))

# Creating plot
plt.boxplot(duration)

# show plot
plt.show()
```



```
data.drop(columns=["pdays", "previous"], axis=1, inplace=True)
```

```
data.describe()
```

	age	duration	campaign
count	32950.000000	32950.000000	32950.000000
mean	40.014112	258.127466	2.560607
std	10.403636	258.975917	2.752326
min	17.000000	0.000000	1.000000
25%	32.000000	103.000000	1.000000
50%	38.000000	180.000000	2.000000
75%	47.000000	319.000000	3.000000
max	98.000000	4918.000000	56.000000

```
# compute interquantile range to calculate the boundaries
```

```
lower_boundries= []
```

```
upper_boundries= []
```

```
for i in ["age", "duration", "campaign"]:
```

```
    IQR= data[i].quantile(0.75) - data[i].quantile(0.25)
```

```
    lower_bound= data[i].quantile(0.25) - (1.5*IQR)
```

```
    upper_bound= data[i].quantile(0.75) + (1.5*IQR)
```

```
    print(i, ":", lower_bound, ",", upper_bound)
```

```
    lower_boundries.append(lower_bound)
```

```
    upper_boundries.append(upper_bound)
```

```
age : 9.5 , 69.5
duration : -221.0 , 643.0
campaign : -2.0 , 6.0
```

```
lower_boundries
```

```
[9.5, -221.0, -2.0]
```

```
upper_boundries
```

```
[69.5, 643.0, 6.0]
```

```
# replace the all the outliers which is greater then upper boundary by upper boundary
```

```
j = 0
for i in ["age", "duration", "campaign"]:
    data.loc[data[i] > upper_boundries[j], i] =
int(upper_boundries[j])
    j = j + 1
```

```
# without outliers
```

```
data.describe()
```

	age	duration	campaign
count	32950.000000	32950.000000	32950.000000
mean	39.929894	234.923915	2.271077
std	10.118566	176.854558	1.546302
min	17.000000	0.000000	1.000000
25%	32.000000	103.000000	1.000000
50%	38.000000	180.000000	2.000000
75%	47.000000	319.000000	3.000000
max	69.000000	643.000000	6.000000

```
age = data["age"]
```


```
fig = plt.figure(figsize =(10, 7))
```

```
# Creating plot
```


```
plt.boxplot(age)
```

```
# show plot
```


```
plt.show()
```




```
campaign = data["campaign"]  
  
fig = plt.figure(figsize =(10, 7))  
  
# Creating plot  
plt.boxplot(campaign)  
  
# show plot  
plt.show()
```





```
duration = data["duration"]  
  
fig = plt.figure(figsize =(10, 7))  
  
# Creating plot  
plt.boxplot(duration)  
  
# show plot  
plt.show()
```



```
# plotting histogram for each numerical variable
plt.style.use("ggplot")
for column in ["age", "duration", "campaign"]:
    plt.figure(figsize=(20,4))
    plt.subplot(121)
    sns.distplot(data[column], kde=True)
    plt.title(column)
```





#categorical features

cat_var

```
Index(['job', 'marital', 'education', 'default', 'housing', 'loan',
      'contact',
      'month', 'day_of_week', 'poutcome', 'y'],
      dtype='object')
```

check categorical class

```
for i in cat_var:
    print(i, ":", data[i].unique())
```

```
job : ['blue-collar' 'entrepreneur' 'retired' 'admin.' 'student'
      'services'
      'technician' 'self-employed' 'management' 'unemployed' 'housemaid']
marital : ['married' 'divorced' 'single']
education : ['basic.9y' 'university.degree' 'basic.4y' 'high.school']
```

```

    'professional.course' 'basic.6y' 'illiterate']
default : ['no' 'yes']
housing : ['no' 'yes']
loan : ['no' 'yes']
contact : ['cellular' 'telephone']
month : ['nov' 'jul' 'may' 'jun' 'aug' 'mar' 'oct' 'apr' 'sep' 'dec']
day_of_week : ['wed' 'mon' 'tue' 'fri' 'thu']
poutcome : ['nonexistent' 'failure' 'success']
y : ['no' 'yes']

```

```

from sklearn.preprocessing import LabelEncoder

```

```

# initializing label encoder

```

```

le= LabelEncoder()

```

```

# iterating through each categorical feature and label encoding them

```

```

for feature in cat_var:
    data[feature]= le.fit_transform(data[feature])

```

```

# label encoded dataset

```

```

data.head()

```

	age	job	marital	education	default	housing	loan	contact
month \								
0	49	1	1	2	0	0	0	0
7								
1	37	2	1	6	0	0	0	1
7								
2	69	5	1	0	0	0	0	0
3								
3	36	0	1	6	0	1	0	1
6								
4	59	5	0	6	0	0	0	0
4								

	day_of_week	duration	campaign	poutcome	y
0	4	227	4	1	0
1	4	202	2	0	0
2	1	643	1	1	1
3	1	120	2	1	0
4	3	368	2	1	0

```

# feature variables

```

```

x= data.iloc[:, :-1]

```

```

# target variable

```

```

y= data.iloc[:, -1]


```

```

plt.figure(figsize=(15,7))
sns.heatmap(data.corr(), annot=True)

```

<matplotlib.axes._subplots.AxesSubplot at 0x7f708aa93710>



#initialising oversampling

```
smote= SMOTETomek(0.75)
```

#implementing oversampling to training data

```
x_sm, y_sm= smote.fit_resample(x,y)
```

x_sm and y_sm are the resampled data

target class count of resampled dataset

```
y_sm.value_counts()
```

```
0    28961
```

```
1    21651
```

```
Name: y, dtype: int64
```

```
x_train, x_test, y_train, y_test= train_test_split(x_sm, y_sm,  
test_size=0.2, random_state=42)
```

selecting the classifier

```
log_reg= LogisticRegression()
```

selecting hyperparameter tuning

```
log_param= {"C": 10.0*np.arange(-2,3), "penalty": ["l1", "l2"]}
```

defining stratified Kfold cross validation

```
cv_log= StratifiedKFold(n_splits=5)
```

using gridsearch for respective parameters

```
gridsearch_log= GridSearchCV(log_reg, log_param, cv=cv_log, scoring=  
"f1_macro", n_jobs=-1, verbose=2)
```

```
# fitting the model on resampled data
```

```
gridsearch_log.fit(x_train, y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=None,
shuffle=False),
              estimator=LogisticRegression(), n_jobs=-1,
              param_grid={'C': array([1.e-02, 1.e-01, 1.e+00, 1.e+01,
1.e+02]),
                          'penalty': ['l1', 'l2']},
              scoring='f1_macro', verbose=2)
```

```
# checking model performance
```

```
y_predicted= gridsearch_log.predict(x_test)
```

```
cm= confusion_matrix(y_test, y_predicted)
```


```
sns.heatmap(cm, annot=True)
```

```
print("Accuracy of Logistic Regression:", accuracy_score(y_test,
y_predicted))
```

```
print(classification_report(y_test, y_predicted))
```

Accuracy of Logistic Regression: 0.797095722611874

	precision	recall	f1-score	support
0	0.81	0.85	0.83	5813
1	0.78	0.73	0.75	4310
accuracy			0.80	10123
macro avg	0.79	0.79	0.79	10123
weighted avg	0.80	0.80	0.80	10123



```
# random forest
rf= RandomForestClassifier()

rf_param= {
    "n_estimators": [int(x) for x in np.linspace(start=100,
stop=1000, num=10)],
    "max_features": ["auto", "sqrt", "log2"],
    # "max_depth": [4,5,6,7,8],
    "max_depth": [int(x) for x in np.linspace(start=5, stop=30,
num=6)],
    "min_samples_split": [5,10,15,100],
    "min_samples_leaf": [1,2,5,10],
    "criterion":['gini', 'entropy']
}

cv_rf= StratifiedKFold(n_splits=5)

randomsearch_rf= RandomizedSearchCV(rf, rf_param, cv=cv_rf, scoring=
"f1_macro", n_jobs=-1, verbose=2, n_iter=10)

randomsearch_rf.fit(x_train, y_train)

Fitting 5 folds for each of 10 candidates, totalling 50 fits

RandomizedSearchCV(cv=StratifiedKFold(n_splits=5, random_state=None,
shuffle=False),
    estimator=RandomForestClassifier(), n_jobs=-1,
    param_distributions={'criterion': ['gini',
'entropy'],
                        'max_depth': [5, 10, 15, 20,
```

```

25, 30],
'sqrt',
10],
15, 100],
300, 400,
700, 800,
'max_features': ['auto',
                  'log2'],
'min_samples_leaf': [1, 2, 5,
                      10],
'min_samples_split': [5, 10,
                       20, 30, 40, 50, 60, 70, 80, 90, 100],
'n_estimators': [100, 200,
                  300, 400, 500, 600,
                  700, 800, 900, 1000]},
scoring='f1_macro', verbose=2)

```

checking model performance


```

y_predicted_rf= randomsearch_rf.predict(x_test)
sns.heatmap(confusion_matrix(y_test, y_predicted_rf), annot=True)
print("Accuracy of Random Forest Model:",accuracy_score(y_test,
y_predicted_rf))
print(classification_report(y_test, y_predicted_rf))

```

Accuracy of Random Forest Model: 0.9015114096611676

	precision	recall	f1-score	support
0	0.93	0.90	0.91	5813
1	0.87	0.91	0.89	4310
accuracy			0.90	10123
macro avg	0.90	0.90	0.90	10123
weighted avg	0.90	0.90	0.90	10123



#Artificial Neural Network

```
ANN_Classifier=Sequential()
ANN_Classifier.add(Dense(units = 6, kernel_initializer = 'uniform',
activation = 'relu', input_dim = 13))
ANN_Classifier.add(Dense(units = 6, kernel_initializer = 'uniform',
activation = 'relu'))
ANN_Classifier.add(Dense(units = 1, kernel_initializer = 'uniform',
activation = 'sigmoid'))
ANN_Classifier.compile(optimizer = 'adam', loss =
'binary_crossentropy', metrics = ['accuracy'])
ANN_Classifier.fit(x_train, y_train, batch_size = 10, epochs = 20)
ANN_Classifier.predict(x_test)
```

Epoch 1/20

4049/4049 [=====] - 12s 3ms/step - loss: 0.4667 - accuracy: 0.7735

Epoch 2/20

4049/4049 [=====] - 11s 3ms/step - loss: 0.4201 - accuracy: 0.8025

Epoch 3/20

4049/4049 [=====] - 12s 3ms/step - loss: 0.4162 - accuracy: 0.8027

Epoch 4/20

4049/4049 [=====] - 11s 3ms/step - loss: 0.4157 - accuracy: 0.8046

Epoch 5/20

4049/4049 [=====] - 11s 3ms/step - loss: 0.4135 - accuracy: 0.8044

Epoch 6/20

4049/4049 [=====] - 11s 3ms/step - loss:

```
0.4121 - accuracy: 0.8046
Epoch 7/20
4049/4049 [=====] - 22s 5ms/step - loss:
0.4115 - accuracy: 0.8041
Epoch 8/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4093 - accuracy: 0.8072
Epoch 9/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4089 - accuracy: 0.8056
Epoch 10/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4072 - accuracy: 0.8074
Epoch 11/20
4049/4049 [=====] - 12s 3ms/step - loss:
0.4060 - accuracy: 0.8065
Epoch 12/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4051 - accuracy: 0.8067
Epoch 13/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4056 - accuracy: 0.8075
Epoch 14/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4047 - accuracy: 0.8071
Epoch 15/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4036 - accuracy: 0.8071
Epoch 16/20
4049/4049 [=====] - 13s 3ms/step - loss:
0.4040 - accuracy: 0.8069
Epoch 17/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4035 - accuracy: 0.8071
Epoch 18/20
4049/4049 [=====] - 10s 3ms/step - loss:
0.4032 - accuracy: 0.8083
Epoch 19/20
4049/4049 [=====] - 11s 3ms/step - loss:
0.4022 - accuracy: 0.8090
Epoch 20/20
4049/4049 [=====] - 10s 3ms/step - loss:
0.4020 - accuracy: 0.8085

array([[0.23497224],
       [0.00383079],
       [0.21346164],
       ...,
       [0.6979944 ]],
```



```

2    55    5        0        5        1        2    0        0
3
3    44    2        1        0        1        0    0        1
4
4    28    0        2        3        0        0    0        0
5

```

```

    day_of_week  duration  campaign  poutcome
0              3      131         5         1
1              3      100         1         1
2              2      131         2         1
3              3       48         2         1
4              0      144         2         1

```

```

x_test = x_test.to_numpy()
x_train = x_train.to_numpy()

```

```

x_train.shape

```

```

(40489, 13)

```

```

x_train = np.reshape(x_train, (x_train.shape[0], 1,
x_train.shape[1]))
x_test = np.reshape(x_test, (x_test.shape[0], 1, x_test.shape[1]))
batch_size = 32

```

```

#LSTM Model

```

```

model = Sequential()
model.add(LSTM(4,input_dim=13))
model.add(Dropout(0.1))
model.add(Dense(1))
model.add(Activation('sigmoid'))

```

```

model.summary()

```

```

Model: "sequential_5"

```

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 4)	288
dropout_1 (Dropout)	(None, 4)	0
dense_13 (Dense)	(None, 1)	5
activation_1 (Activation)	(None, 1)	0

```

=====
Total params: 293
Trainable params: 293
Non-trainable params: 0

```

```

model.compile(loss='binary_crossentropy',optimizer='adam',metrics=['accuracy'])
checkpointer = callbacks.ModelCheckpoint(filepath="checkpoint-
{epoch:02d}.hdf5", verbose=1, save_best_only=True,
monitor='val_acc',mode='max')
model.fit(x_train, y_train, batch_size=batch_size, epochs=10,
validation_data=(x_test, y_test))

```

Epoch 1/10

```

1266/1266 [=====] - 10s 5ms/step - loss:
0.6076 - accuracy: 0.6864 - val_loss: 0.5762 - val_accuracy: 0.7382

```

Epoch 2/10

```

1266/1266 [=====] - 5s 4ms/step - loss:
0.5486 - accuracy: 0.7322 - val_loss: 0.5092 - val_accuracy: 0.7435

```

Epoch 3/10

```

1266/1266 [=====] - 5s 4ms/step - loss:
0.4973 - accuracy: 0.7574 - val_loss: 0.4687 - val_accuracy: 0.7746

```

Epoch 4/10

```

1266/1266 [=====] - 5s 4ms/step - loss:
0.4753 - accuracy: 0.7766 - val_loss: 0.4476 - val_accuracy: 0.7915

```

Epoch 5/10

```

1266/1266 [=====] - 6s 4ms/step - loss:
0.4681 - accuracy: 0.7765 - val_loss: 0.4479 - val_accuracy: 0.7837

```

Epoch 6/10

```

1266/1266 [=====] - 6s 4ms/step - loss:
0.4651 - accuracy: 0.7744 - val_loss: 0.4394 - val_accuracy: 0.7781

```

Epoch 7/10

```

1266/1266 [=====] - 6s 5ms/step - loss:
0.4609 - accuracy: 0.7628 - val_loss: 0.4372 - val_accuracy: 0.7936

```

Epoch 8/10

```

1266/1266 [=====] - 6s 5ms/step - loss:
0.4561 - accuracy: 0.7593 - val_loss: 0.4312 - val_accuracy: 0.8000

```

Epoch 9/10

```

1266/1266 [=====] - 5s 4ms/step - loss:
0.4555 - accuracy: 0.7593 - val_loss: 0.4317 - val_accuracy: 0.7975

```

Epoch 10/10

```

1266/1266 [=====] - 5s 4ms/step - loss:
0.4540 - accuracy: 0.7637 - val_loss: 0.4300 - val_accuracy: 0.7998

```

<keras.callbacks.History at 0x7f708bfc6f50>

```

loss, accuracy = model.evaluate(x_test, y_test)
print("Accuracy of LSTM Model: %.2f%%" % (accuracy*100))
y_pred = model.predict(x_test)
y_pred = (y_pred > 0.5)
confusion_matrix(y_test, y_pred)
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True)
print(classification_report(y_test, y_pred))

```


```

317/317 [=====] - 1s 3ms/step - loss: 0.4300
- accuracy: 0.7998

```

Accuracy of LSTM Model: 79.98%

	precision	recall	f1-score	support
0	0.84	0.81	0.82	5813
1	0.75	0.79	0.77	4310
accuracy			0.80	10123
macro avg	0.80	0.80	0.80	10123
weighted avg	0.80	0.80	0.80	10123



```
training_set, test_set = train_test_split(data, test_size = 0.2,  
random_state = 1)
```

```
x_train = training_set.iloc[:,0:13].values
```

```
y_train = training_set.iloc[:,13].values
```

```
x_test = test_set.iloc[:,0:13].values
```

```
y_test = test_set.iloc[:,13].values
```

```
#SVM Model
```

```
classifier = SVC(kernel='rbf', random_state = 1)
```

```
classifier.fit(x_train,y_train)
```

```
SVC(random_state=1)
```

```
y_pred = classifier.predict(x_test)
```

```
test_set["Predictions"] = y_pred
```

```
from sklearn.metrics import confusion_matrix
```

```
cm = confusion_matrix(y_test,y_pred)
```


```

accuracy = float(cm.diagonal().sum())/len(y_test)
print("\nAccuracy Of SVM For The Given Dataset : ", accuracy)
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True)
print(classification_report(y_test, y_pred))

```

Accuracy Of SVM For The Given Dataset : 0.887556904400607

	precision	recall	f1-score	support
0	0.89	1.00	0.94	5849
1	0.00	0.00	0.00	741
accuracy			0.89	6590
macro avg	0.44	0.50	0.47	6590
weighted avg	0.79	0.89	0.83	6590



```

# predicting the test data
y_predicted= randomsearch_rf.predict(test_data)
y_predicted

array([0, 0, 0, ..., 1, 0, 0])

# dataset of predicted values for target variable y
prediction= pd.DataFrame(y_predicted, columns=["y_predicted"])
prediction_dataset= pd.concat([test_data, prediction], axis=1)
prediction_dataset['y_predicted'] =
prediction_dataset['y_predicted'].map({1: 'yes', 0: 'no'})
prediction_dataset.to_csv('Output.csv')
prediction_dataset

```

	month	age \	job	marital	education	default	housing	loan	contact
0	3	32	4	0	6	0	0	0	0
1	4	37	10	3	6	0	0	0	0
2	3	55	5	0	5	1	2	0	0
3	4	44	2	1	0	1	0	0	1
4	5	28	0	2	3	0	0	0	0
...
8233	6	48	4	1	2	0	2	0	0
8234	6	30	7	2	3	0	2	0	0
8235	4	33	7	1	3	0	0	0	0
8236	6	44	1	1	1	0	2	2	1
8237	6	42	1	1	2	1	2	0	0

	day_of_week	duration	campaign	poutcome	y_predicted
0	3	131	5	1	no
1	3	100	1	1	no
2	2	131	2	1	no
3	3	48	2	1	no
4	0	144	2	1	no
...
8233	3	554	1	1	yes
8234	0	159	1	1	no
8235	1	472	1	0	yes
8236	1	554	5	1	no
8237	3	83	5	1	no

[8238 rows x 14 columns]